

Ad Copilot — Release Plan

Status: Draft v1.0 **Owner:** Product, in partnership with Engineering & Design **Related docs:** Case Study Narrative · PRD · User Stories · AI Workflow Design · Metrics & Attribution Spec · Architecture Doc

This is the final doc in the set. It operationalizes the PRD's three-phase scope (§11) into concrete entry/exit gates, rollback criteria, and a beta strategy — the release shouldn't be a launch date, it should be a sequence of trust-earning gates, consistent with the product's core thesis from the Case Study.

1. Release Philosophy

Ad Copilot's biggest adoption risk isn't feature completeness, it's a marketer's trust in AI-driven action on their live budget. The release plan mirrors that: **each phase expands scope only after the previous phase has demonstrated the trust and quality signals it needs to**. No phase is gated purely on a calendar date.

2. Phase 1 — MVP: Creative Generation + Suggest-Lane Review (Meta only)

Scope: US-1, US-2, US-3, US-5, US-6, US-7, US-8 (Suggest/Escalate only, no Auto-apply), US-9, US-10, US-12, US-13 (per PRD §11, User Stories P0 set).

Entry criteria:

- Meta Ads API integration complete and passing data-accuracy validation against a manual spot-check of 10 real (beta partner) accounts.
- Brand-safety filter validated against a test set with zero false negatives on a curated list of known-bad terms.
- Cost-cap enforcement (US-12) load-tested to confirm it blocks generation reliably at the boundary, not just on average.

Beta cohort: 5-10 design-partner accounts, hand-selected for a mix of e-commerce (ROAS-target) and lead-gen (CPA-target) use cases, actively managing \$50K+/month so drift signals have enough data to be meaningful (ties to Metrics & Attribution Spec §4's

data-volume thresholds).

Exit criteria (to proceed to Phase 2):

- Suggestion acceptance rate $\geq 40\%$ across the beta cohort, sustained over 4+ weeks (PRD goal, §3).
- Zero incidents of brand-safety filter failure reaching a live account.
- No unresolved P0/P1 bugs in the Action Executor or Audit Log (these are the trust-critical paths).
- Qualitative: beta partners report the recommendation *reasoning* (not just the recommendation) is useful — this is checked via structured beta interviews, not just usage analytics, since acceptance rate alone can mask a “rubber-stamping” failure mode.

Rollback/kill-switch: a per-account “pause AI” toggle exists from day one (not phased in later) — any beta partner can fully disable generation and recommendations instantly, reverting to platform-native behavior with zero residual effect.

3. Phase 2 — Multi-Platform + Budget/Bid Recommendations

Scope: Add Google Ads integration; extend recommendation engine from creative-only to budget/bid recommendations (still Suggest-lane only, no Auto-apply); US-4, US-11, US-14 (P1 stories) considered for inclusion based on Phase 1 learnings.

Entry criteria:

- Phase 1 exit criteria met and sustained (not a one-time snapshot).
- Google Ads developer token approval secured (Architecture Doc §3.2 flagged this as a lead-time risk — it must be resolved before this phase is scheduled, not during it).
- Confidence-scoring calibration reviewed against real Phase 1 outcome data (AI Workflow Design §7) before extending the same scoring approach to a new action type (budget/bid changes carry more direct financial risk than creative swaps).

Rollout approach: budget/bid recommendations ship behind a feature flag, enabled first only for the Phase 1 beta cohort (who already have trust and history with the product) before any new-account rollout.

Exit criteria (to proceed to Phase 3):

- Budget/bid recommendation acceptance rate reaches parity with creative-recommendation acceptance rate (a lower rate here would suggest marketers trust

budget suggestions less, which must be understood before granting any autonomy in Phase 3).

- Incremental ROAS north-star metric (Metrics & Attribution Spec §6) shows a positive, statistically-reasonable trend across the expanded cohort, not just the original beta accounts.
- No Escalate-lane item has gone unreviewed past its (to-be-defined) staleness threshold during the phase, validating the human-review flow scales past the original small cohort.

4. Phase 3 — Auto-Apply Lane

Scope: US-8's Auto-apply lane goes live; US-10's rollback mechanism is now load-bearing in production, not just a Phase 1/2 safety net that was never triggered.

Entry criteria:

- An account only becomes Auto-apply-eligible after meeting the trust-tier bar defined in AI Workflow Design §5 (sufficient data volume + sustained acceptance history) — this is enforced per-account, not toggled globally.
- Platform-level hard ceilings on Auto-apply blast radius (AI Workflow Design §5) are implemented and confirmed non-configurable by admins, closing the loop on a risk flagged back in the AI Workflow Design doc.
- Rollback mechanism tested under real failure injection (e.g., simulated bad recommendation) in a staging environment before any real account is exposed to it.

Rollout approach: opt-in only. Even trust-tier-eligible accounts must explicitly enable Auto-apply per campaign — this is never turned on by default, reinforcing the product's core "earned autonomy" positioning rather than presenting it as an automatic upgrade.

Success criteria: Auto-apply accounts show equal or better incremental ROAS than their own Suggest-lane history for the same campaigns, and rollback is invoked at a low, stable rate (a rising rollback rate is treated as a signal to tighten thresholds, not a normal cost of automation).

5. Feature Flags

Flag	Controls	Default
<code>platform.google_ads</code>	Google Ads integration availability	Off until Phase 2 entry criteria met

<code>recs.budget_bid</code>	Budget/bid recommendation generation (vs. creative-only)	Off until Phase 2
<code>lane.auto_apply</code>	Whether Auto-apply lane can ever be reached for an account	Off until Phase 3, then opt-in per account even when globally enabled
<code>account.ai_paused</code>	Per-account master kill-switch	Available from Phase 1 launch, always overridable by the account owner

6. Post-Launch Monitoring (all phases)

- Weekly review of the guardrail metrics from the Metrics & Attribution Spec (§7): LLM cost as % of managed spend, acceptance rate, escalation rate.
- Any single-account cost-cap breach investigated same-week, not batched into a monthly review, since cost runaway is a fast-moving risk category.
- Confidence calibration review cadence formalized (tie-in to AI Workflow Design §7) — this doc treats it as an ongoing release-health activity, not a one-time launch task.

7. What “Done” Looks Like for This Portfolio Series

Taken together, the six documents in this set — Case Study, PRD, User Stories, AI Workflow Design, Metrics & Attribution Spec, Architecture Doc, and this Release Plan — trace a single consistent product decision (earn autonomy incrementally, make the AI’s boundaries explicit) through every layer: narrative, requirements, backlog, AI behavior, marketing-metric grounding, technical architecture, and rollout sequencing.