

Ad Copilot — User Story Backlog & Acceptance Criteria

Status: Draft v1.0 **Related docs:** Case Study Narrative · PRD · AI Workflow Design (next) · Metrics & Attribution Spec (next) · Architecture Doc (next) · Release Plan (next) **Format:** Stories are grouped by epic, mapped to PRD Functional Requirements (FR-#), with Gherkin-style acceptance criteria. Points are rough t-shirt-to-Fibonacci estimates for planning illustration, not a claim of actual engineering estimation.

Epic 1: Account Connection & Setup

US-1 — Connect ad account

As a performance marketer, **I want to** connect my Meta Ads account via OAuth **so that** Ad Copilot can read my campaign data. *Maps to: onboarding flow, PRD §8 flow 1* **Priority:** P0 · **Points:** 5

- Given I am on the onboarding screen, when I click “Connect Meta Ads,” then I am redirected to Meta’s OAuth consent screen.
- Given I grant permission, when I return to Ad Copilot, then my last 90 days of campaign, ad set, and ad-level performance data begins syncing.
- Given the sync is in progress, when I view the dashboard, then I see a sync-status indicator rather than an empty or misleading state.
- Given I deny permission or the OAuth flow fails, when I return to Ad Copilot, then I see a clear error message and a retry option.

US-2 — Set performance targets per campaign

As a performance marketer, **I want to** set a target ROAS or CPA per campaign **so that** the system knows when performance has drifted. *Maps to: FR-6, PRD §8 flow 2* **Priority:** P0 · **Points:** 3

- Given I’ve connected an account, when I open a campaign’s settings, then I can input a target ROAS or target CPA (not both, to avoid conflicting signals).
- Given no target is set, when drift monitoring runs, then that campaign is excluded from

recommendations and is visibly marked “no target set” rather than silently skipped.

Epic 2: Creative Generation

US-3 — Generate ad copy variants on demand

As a performance marketer, **I want to** generate new ad copy variants for a specific ad set **so that** I can replace fatiguing creative quickly. *Maps to: FR-1, FR-3* **Priority:** P0 · **Points:** 8

- Given I select an ad set and click “Generate variants,” when generation completes, then I see 3-5 headline/primary-text/CTA combinations.
- Given variants are generated, when I view them, then each displays a confidence score.
- Given the product feed for that ad set is missing required fields (e.g., no product title), when I attempt generation, then I receive a specific error naming the missing field, not a generic failure.

US-4 — Scheduled creative refresh

As a performance marketer, **I want to** schedule automatic variant generation for ad sets that show fatigue signals **so that** I don’t have to manually monitor for decay. *Maps to: FR-3* **Priority:** P1 · **Points:** 5

- Given I enable “auto-refresh” on an ad set, when the system detects a fatigue signal (e.g., CTR decline over a defined window), then new variants are generated and placed in the review queue.
- Given auto-refresh is enabled, when no fatigue signal occurs, then no unnecessary generation calls are made (cost control tie-in, see FR-13).

US-5 — Brand-safety filtering on generated copy

As a brand/compliance stakeholder, **I want** generated ad copy to be screened against a wordlist **so that** non-compliant or off-brand copy never reaches a reviewer. *Maps to: FR-4* **Priority:** P0 · **Points:** 5

- Given a generated variant contains a blocked term, when generation completes, then that variant is suppressed and never appears in the review queue.
 - Given a variant is suppressed, when an admin views generation logs, then the suppressed variant and the matched term are visible for audit purposes (but not shown to the general reviewer to avoid noise).
-

Epic 3: Monitoring & Recommendations

US-6 — Detect performance drift

As the system, **I want to** compare live performance against campaign targets on an hourly basis **so that** drift is caught before days of wasted spend. *Maps to: FR-5, FR-6* **Priority:** P0 · **Points:** 8

- Given a campaign has a target CPA, when actual CPA exceeds target by more than the configured threshold for the configured duration, then a recommendation is generated.
- Given the ad platform API is rate-limited or returns stale data, when the sync job runs, then the system retries with backoff and flags the data as stale in the UI rather than acting on incomplete data.

US-7 — Receive a budget/bid recommendation with reasoning

As a performance marketer, **I want to** see why a recommendation was made **so that** I can evaluate it instead of blindly trusting or distrusting it. *Maps to: FR-7* **Priority:** P0 · **Points:** 5

- Given a recommendation is generated, when I open it, then I see the specific metric that triggered it, the current vs. target value, and the proposed action.
 - Given a recommendation is generated, when I open it, then I see a confidence score and a plain-language explanation, not just a raw model output.
-

Epic 4: Review Lanes & Human-in-the-Loop

US-8 — Lane assignment for AI actions

As the system, **I want to** classify every AI action into Auto-apply, Suggest, or Escalate **so that** risk is proportional to autonomy. *Maps to: FR-8, FR-9* **Priority:** P0 · **Points:** 8

- Given an action's confidence score and blast radius are calculated, when lane assignment runs, then the action is placed in exactly one lane per the configured rules.
- Given an account is new (below a configured data-volume threshold), when any action is generated, then it is placed in Suggest or Escalate regardless of confidence score (new accounts cannot land in Auto-apply).

US-9 — Review and act on a queued recommendation

As a performance marketer, **I want to** approve, reject, or edit a queued recommendation **so**

that I retain control over budget-impacting actions. *Maps to: FR-10* **Priority:** P0 · **Points:** 5

- Given a recommendation is in the Suggest lane, when I click "Approve," then the action is executed against the ad platform API and logged.
- Given a recommendation is in the Suggest lane, when I click "Edit," then I can modify the proposed value (e.g., change -15% budget to -10%) before approving.
- Given I reject a recommendation, when I do so, then I can optionally provide a reason, which is stored to improve future confidence scoring.

US-10 — Audit log and rollback

As a performance marketer, **I want to** see a full history of AI actions, including auto-applied ones, and roll them back **so that** I never lose visibility or control. *Maps to: FR-11* **Priority:** P0 · **Points:** 5

- Given any AI action has been taken (any lane), when I view the audit log, then I see the action, timestamp, lane, confidence score, and current status.
- Given an Auto-apply action was taken within the last 5 minutes, when I click "Rollback," then the prior budget/bid value is restored and this is logged as a distinct event.

US-11 — Configure lane thresholds per account

As an admin user, **I want to** adjust the confidence/blast-radius thresholds for lane assignment **so that** the system's risk tolerance matches our account's needs. *Maps to: FR-12* **Priority:** P1 · **Points:** 3

- Given I am an admin, when I adjust the Auto-apply confidence threshold upward, then fewer actions qualify for Auto-apply going forward (existing queued items are unaffected).
- Given I attempt to set a threshold that would allow unbounded budget risk (e.g., no blast-radius cap), when I save, then the system blocks the change and explains why.

Epic 5: Cost Controls

US-12 — Enforce monthly generation caps

As the business, **I want to** cap LLM generation calls per account per month **so that** AI cost doesn't scale unpredictably against a customer's plan tier. *Maps to: FR-13* **Priority:** P0 · **Points:** 5

- Given an account has reached its monthly generation cap, when a user or scheduled job attempts further generation, then the request is blocked with a clear message and an option to upgrade plan tier.
- Given an account is approaching 80% of its cap, when this threshold is crossed, then the account owner receives a proactive notification.

US-13 — Tiered model routing

As the system, **I want to** use a lower-cost model for first-pass generation and reserve the higher-cost model for top candidates **so that** generation cost per account stays proportionate. *Maps to: FR-14* **Priority:** P0 · **Points:** 8

- Given a generation request produces an initial candidate set, when the top-N candidates are selected by predicted score, then only those N are refined using the higher-cost model.
- Given the refinement step fails or times out, when this occurs, then the first-pass candidates are still surfaced to the user rather than the whole generation failing silently.

US-14 — Cache reuse for unchanged inputs

As the system, **I want to** reuse prior generation outputs when the underlying product-feed input hasn't changed **so that** we avoid redundant, wasteful model calls. *Maps to: FR-15* **Priority:** P1 · **Points:** 5

- Given a product feed entry is unchanged since the last generation within the caching window, when a new generation request is made for that entry, then cached outputs are served instead of a new model call.
- Given cached outputs are served, when the user views them, then they are visually indistinguishable in quality expectations from freshly generated ones (no "stale" stigma in UI).

Backlog Summary (for sprint planning reference)

Epic	Stories	P0 count	Total points (illustrative)
Account Connection & Setup	2	2	8
Creative Generation	3	2	18
Monitoring & Recommendations	2	2	13

Review Lanes & Human-in-the-Loop	4	3	21
Cost Controls	3	2	18

Suggested MVP cut: All P0 stories (13 stories, ~72 points), deferring US-4, US-11, and US-14 (P1) to Phase 2 per the PRD's phased scope.